

Ad Deduplication with images

Description

Schibsted owns several online marketplaces such as e.g. *finn.no*. In these marketplaces users can post classified ads for objects they wish to sell. For various reasons, users might post multiple classified ads for the same item, i.e. creating duplicate classified ads. We'd like to avoid this, as it makes the user experience more confusing for other users.

In this task, we ask that you write code to detect such duplicate classified ads just by comparing the *images* from the classified ads. For simplicity we assume that each classified ad only has a single image. We've included some examples of images that might be from multiple such duplicate ads.

The task formulation is a bit vague on purpose, enabling you to make your own choices where needed.

Instructions

- Download examples of images from duplicate ads from the link https://drive.google.com/open?id=1GqnKQPjb_4x5pL8Wdlylr5fqfMsg-GCu
- For your reference, examples of images from non-duplicate ads can be downloaded from the link https://drive.google.com/open?id=1Ok-yU50_bzvCedRha6XwznpO621RiX-z
- Your task is to create code that takes as input two images and returns some value indicating whether they are duplicates or not
- Please submit a running piece of code, documented as you please (in the code, separate docs, document by tests, whatever makes the code understandable)
- We are interested in seeing your ideas, and your coding style/skills. Don't worry if the results are sub-optimal
- If you have ideas on how to improve the solution further, please write that down as well
- Spend approximately 3 hours on the task